

# ML Prediction writeup

*Stefaan Delanghe*

*9/30/2016*

## Assignment

Using devices such as Jawbone Up, Nike FuelBand, and Fitbit it is now possible to collect a large amount of data about personal activity relatively inexpensively. These type of devices are part of the quantified self movement ??? a group of enthusiasts who take measurements about themselves regularly to improve their health, to find patterns in their behavior, or because they are tech geeks. One thing that people regularly do is quantify how much of a particular activity they do, but they rarely quantify how well they do it.

In this project, your goal will be to use data from accelerometers on the belt, forearm, arm, and dumbbell of 6 participants. They were asked to perform barbell lifts correctly and incorrectly in 5 different ways. More information is available from the website here: <http://groupware.les.inf.puc-rio.br/har> (see the section on the Weight Lifting Exercise Dataset).

## Preparation

Following libraries are required for this assignment.

```
library(dplyr)
library(ggplot2)
library(caret)
library(rpart)
library(rpart.plot)
```

## Load data

We download the prediction writeup data and load the csv files.

```
download.file("https://d396qusza40orc.cloudfront.net/predmachlearn/pml-training.csv",
  destfile = "pml-training.csv")
download.file("https://d396qusza40orc.cloudfront.net/predmachlearn/pml-testing.csv",
  destfile = "pml-testing.csv")
pml.training <- read.csv("pml-training.csv")
pml.testing <- read.csv("pml-testing.csv")
```

## Tidy data

Before we get into the prediction we require to clean the data. In this case the data contains alot of NA values, which can influence our prediction in a negative way. So we will clean up the data in order to have an accurate prediction as possible.

Our first clean up is reducing the number of columns to the ones that have been mentioned in the assignment. We only use the columns that have belt, forearm, arm, and dumbbell.

We only use one dataset in order to perform the clean up. We bind the training and test set into the pml set. After that we filter only the columns with arm, beld, dumbell, classe.

The training and test set do not have the same number of columns. So they can not be bind by default. For this we will add each missing column to each set.

```
pml.testing$classe <- NA
pml.training$problem_id <- NA

pml <- rbind(pml.training, pml.testing)
pml <- select(pml, matches("arm|belt|dumbell|classe|problem_id"))
```

In order for our prediction to be most accurate we will remove the rows that contain NA values.

```
pml.cols <- c(colnames(pml[colSums(is.na(pml)) == 0]), colnames(pml[115:116]))
pml.feats <- pml[pml.cols]
pml.feats %>% filter(complete.cases(.))
```

```
## [1] roll_belt          pitch_belt          yaw_belt
## [4] total_accel_belt   gyros_belt_x        gyros_belt_y
## [7] gyros_belt_z       accel_belt_x        accel_belt_y
## [10] accel_belt_z       magnet_belt_x       magnet_belt_y
## [13] magnet_belt_z      roll_arm            pitch_arm
## [16] yaw_arm            total_accel_arm     gyros_arm_x
## [19] gyros_arm_y        gyros_arm_z         accel_arm_x
## [22] accel_arm_y        accel_arm_z         magnet_arm_x
## [25] magnet_arm_y       magnet_arm_z        roll_forearm
## [28] pitch_forearm      yaw_forearm         total_accel_forearm
## [31] gyros_forearm_x    gyros_forearm_y     gyros_forearm_z
## [34] accel_forearm_x    accel_forearm_y     accel_forearm_z
## [37] magnet_forearm_x   magnet_forearm_y     magnet_forearm_z
## [40] classe             problem_id
## <0 rows> (or 0-length row.names)
```

## Train model

Before we train our model we require to the creation of a training and a test set. We use the pml set for the creation of both.

```
pml.part <- createDataPartition(y = pml.feats$classe, p=0.7, list=FALSE)
training <- pml.feats[pml.part,]
testing <- pml.feats[-pml.part,]
```

Now we will train our model with the **random forest** method. This algorithm as it selects most important variables automatically. A **5 fold cross validation** to the training of the model. In order to determine the best method we could have applied an ada boost method on the set. Due to the limitations of the project random forest will suffice.

The training of the model can take a long time. Our computer works usually with one core. We increase the number of cores assigned to the training of the model to increase performance.

```
library(doMC)
registerDoMC(cores = 8)
pml.rf <- train(as.factor(classe)~., data = training[1:40],
               method = "rf", na.action = na.omit,
               trControl = trainControl(method = "cv", 5), allowParallel = TRUE)
print(pml.rf)
```

```
## Random Forest
##
## 13737 samples
## 39 predictor
## 5 classes: 'A', 'B', 'C', 'D', 'E'
##
## No pre-processing
## Resampling: Cross-Validated (5 fold)
## Summary of sample sizes: 10990, 10990, 10987, 10991, 10990
## Resampling results across tuning parameters:
##
## mtry Accuracy Kappa
## 2 0.9881347 0.9849891
## 20 0.9876978 0.9844376
## 39 0.9803448 0.9751346
##
## Accuracy was used to select the optimal model using the largest value.
## The final value used for the model was mtry = 2.
```

This method seems like a good fit as we achieve a high accuracy.

## Confusion and statistics

Now we will make **prediction**, **confusion matrix** that applies to the **20 test cases** available within the testing set.

```
pml.validate <- predict(pml.rf, testing)
confusionMatrix(pml.validate, testing$classe)
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  A    B    C    D    E
##      A 1672    6    1    1    0
##      B   1 1133   14    0    0
##      C   0    0 1004    6    0
##      D   1    0    7  953    0
##      E   0    0    0    4 1082
##
## Overall Statistics
##
##               Accuracy : 0.993
##               95% CI : (0.9906, 0.995)
##      No Information Rate : 0.2845
##      P-Value [Acc > NIR] : < 2.2e-16
##
##               Kappa : 0.9912
##      McNemar's Test P-Value : NA
##
## Statistics by Class:
##
##               Class: A Class: B Class: C Class: D Class: E
```

## Sensitivity	0.9988	0.9947	0.9786	0.9886	1.0000
## Specificity	0.9981	0.9968	0.9988	0.9984	0.9992
## Pos Pred Value	0.9952	0.9869	0.9941	0.9917	0.9963
## Neg Pred Value	0.9995	0.9987	0.9955	0.9978	1.0000
## Prevalence	0.2845	0.1935	0.1743	0.1638	0.1839
## Detection Rate	0.2841	0.1925	0.1706	0.1619	0.1839
## Detection Prevalence	0.2855	0.1951	0.1716	0.1633	0.1845
## Balanced Accuracy	0.9985	0.9958	0.9887	0.9935	0.9996

## Random forest visualization

Here follows a visualization of the partial random forest method applied.

```
pml.tree <- rpart(classe ~ ., data=training, method="class")
prp(pml.tree)
```

